

## ✓ MADHUR GUPTA

**AIM - To fine tune a pre-trained CNN architecture and evaluate its performance on a dataset.**

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
import matplotlib.pyplot as plt
import numpy as np

# Load and preprocess the dataset
# Replace with your actual dataset loading and preprocessing
(x_train, y_train), (x_test, y_test) = keras.datasets.cifar10.load_data()
x_train = x_train.astype("float32") / 255.0
x_test = x_test.astype("float32") / 255.0
y_train = keras.utils.to_categorical(y_train, num_classes=10)
y_test = keras.utils.to_categorical(y_test, num_classes=10)

# Define the pre-trained CNN model
base_model = keras.applications.ResNet50(
    weights="imagenet", include_top=False, input_shape=(32, 32, 3)
)

# Freeze the pre-trained layers
base_model.trainable = False

# Add custom classification layers
inputs = keras.Input(shape=(32, 32, 3))
x = base_model(inputs, training=False)
x = layers.GlobalAveragePooling2D()(x)
x = layers.Dropout(0.2)(x)
outputs = layers.Dense(10, activation="softmax")(x)
model = keras.Model(inputs, outputs)

# Compile the model
model.compile(
    optimizer="adam", loss="categorical_crossentropy", metrics=["accuracy"]
)

# Train the model
epochs = 10 # Adjust as needed
history = model.fit(x_train, y_train, epochs=epochs, validation_data=(x_test, y_test))

# Fine-tune the top layers of the pre-trained model
base_model.trainable = True
for layer in base_model.layers[:100]: # Fine-tune a specific number of layers
    layer.trainable = False

model.compile(
    optimizer=keras.optimizers.Adam(1e-5), # Low learning rate for fine-tuning
    loss="categorical_crossentropy",
    metrics=["accuracy"],
)

fine_tune_epochs = 10 # Adjust as needed
total_epochs = epochs + fine_tune_epochs
history_fine = model.fit(x_train, y_train, epochs=total_epochs, initial_epoch=epochs, validation_data=(x_test, y_test))


# Evaluate the model
loss, accuracy = model.evaluate(x_test, y_test)
print(f"Test Loss: {loss}")
print(f"Test Accuracy: {accuracy}")

#Plot training history
acc = history.history['accuracy'] + history_fine.history['accuracy']
val_acc = history.history['val_accuracy'] + history_fine.history['val_accuracy']
loss = history.history['loss'] + history_fine.history['loss']
val_loss = history.history['val_loss'] + history_fine.history['val_loss']


epochs_range = range(total_epochs)
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>  
 170498071/170498071 — 4s 0us/step  
 Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet50\\_weights\\_tf\\_dim\\_orderi\\_94765736/94765736](https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet50_weights_tf_dim_orderi_94765736/94765736) — 1s 0us/step  
 Epoch 1/10  
 1563/1563 — 42s 18ms/step — accuracy: 0.1695 — loss: 2.3354 — val\_accuracy: 0.2494 — val\_loss: 2.0318  
 Epoch 2/10  
 1563/1563 — 22s 9ms/step — accuracy: 0.2568 — loss: 2.0350 — val\_accuracy: 0.3299 — val\_loss: 1.9075  
 Epoch 3/10  
 1563/1563 — 14s 9ms/step — accuracy: 0.2724 — loss: 2.0110 — val\_accuracy: 0.3305 — val\_loss: 1.8925  
 Epoch 4/10  
 1563/1563 — 15s 9ms/step — accuracy: 0.2766 — loss: 1.9968 — val\_accuracy: 0.3018 — val\_loss: 1.9073  
 Epoch 5/10  
 1563/1563 — 23s 11ms/step — accuracy: 0.2799 — loss: 1.9830 — val\_accuracy: 0.3236 — val\_loss: 1.8769  
 Epoch 6/10  
 1563/1563 — 18s 9ms/step — accuracy: 0.2836 — loss: 1.9831 — val\_accuracy: 0.3327 — val\_loss: 1.8841  
 Epoch 7/10  
 1563/1563 — 20s 9ms/step — accuracy: 0.2848 — loss: 1.9760 — val\_accuracy: 0.3310 — val\_loss: 1.8721  
 Epoch 8/10  
 1563/1563 — 15s 9ms/step — accuracy: 0.2834 — loss: 1.9782 — val\_accuracy: 0.3115 — val\_loss: 1.8879  
 Epoch 9/10  
 1563/1563 — 15s 9ms/step — accuracy: 0.2875 — loss: 1.9748 — val\_accuracy: 0.3297 — val\_loss: 1.8637  
 Epoch 10/10  
 1563/1563 — 15s 9ms/step — accuracy: 0.2846 — loss: 1.9802 — val\_accuracy: 0.3324 — val\_loss: 1.8630  
 Epoch 11/20  
 1563/1563 — 87s 35ms/step — accuracy: 0.1380 — loss: 7.5832 — val\_accuracy: 0.1890 — val\_loss: 2.8255  
 Epoch 12/20  
 1563/1563 — 34s 22ms/step — accuracy: 0.2098 — loss: 2.3339 — val\_accuracy: 0.2880 — val\_loss: 2.1933  
 Epoch 13/20  
 1563/1563 — 40s 21ms/step — accuracy: 0.2756 — loss: 2.0669 — val\_accuracy: 0.3434 — val\_loss: 1.9389  
 Epoch 14/20  
 1563/1563 — 34s 22ms/step — accuracy: 0.3400 — loss: 1.8901 — val\_accuracy: 0.3703 — val\_loss: 1.8178  
 Epoch 15/20  
 1563/1563 — 40s 21ms/step — accuracy: 0.3838 — loss: 1.7544 — val\_accuracy: 0.4111 — val\_loss: 1.7029  
 Epoch 16/20  
 1563/1563 — 33s 21ms/step — accuracy: 0.4248 — loss: 1.6395 — val\_accuracy: 0.4442 — val\_loss: 1.5532  
 Epoch 17/20  
 1563/1563 — 41s 21ms/step — accuracy: 0.4579 — loss: 1.5554 — val\_accuracy: 0.4689 — val\_loss: 1.5216  
 Epoch 18/20  
 1563/1563 — 33s 21ms/step — accuracy: 0.4805 — loss: 1.4817 — val\_accuracy: 0.4653 — val\_loss: 1.5058  
 Epoch 19/20  
 1563/1563 — 41s 21ms/step — accuracy: 0.5037 — loss: 1.4207 — val\_accuracy: 0.5098 — val\_loss: 1.4029  
 Epoch 20/20  
 1563/1563 — 42s 21ms/step — accuracy: 0.5151 — loss: 1.3738 — val\_accuracy: 0.4901 — val\_loss: 1.4461  
 313/313 — 3s 8ms/step — accuracy: 0.4893 — loss: 1.4455  
 Test Loss: 1.446104884147644  
 Test Accuracy: 0.4900999963283539

Training and Validation Accuracy



Training and Validation Loss



```

plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(epochs_range, acc, label='Training Accuracy')


```

```
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')


plt.subplot(1, 2, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()
```



Training and Validation Accuracy



Training and Validation Loss



Start coding or [generate](#) with AI.